

A DISTRIBUTION PLAYBOOK · DIGITAL B2B · SHOPIFY

Getting DynoWeb Its First Paying Customers

A sequenced 90-day plan — interview first, validate the message, then pitch a named list by hand and turn one channel into a repeatable system.

STARTING POINT	THE GOAL	TIME BUDGET
0 real paying customers ~18 real free-tier stores installs stalled to ~1/wk	Target: 50 paying / 60 days Honest ceiling: ~8–15 (see funnel math, §7)	15–20 hrs/week founder-led, hands-on outbound + interviews

“Poor distribution — not product — is the number one cause of startup failure.” In B2B it is truer still: the best tool loses to the one the buyer’s peers already trust. This plan is built on that fact — credibility and proof come before reach.

Prepared 2026-07-17 · grounded in live DynoWeb production data and peer feedback (Charlie Ngo).

The plan, in one breath

Nine steps, in order. The spine is simple: **you cannot write the message, shoot the video, or pitch the list until you have actually listened to merchants.** Everything upstream of the interviews is preparation; everything downstream is execution. The single discipline that keeps it honest is the funnel math in §7 — revisit it whenever a step tempts you to do more of everything instead of more of what works.

STEP	WHAT	WHY IT COMES HERE
1	Define the ICP	You can’t pick a channel or a list before you know exactly who you’re for.
2	Interview — warm 18 first	The customers hand you the words. Do this <i>before</i> any asset work.
3	Pick ≤ 3 value points	Charlie’s note: no focused value prop today. The interviews decide the 3.
4	Remake video, images, site	Only now — built around validated words, not guesses.
5	Bullseye traction research	Map all 19 channels + study how App Store competitors get installs.
6	Test 3 channels, pick 1	The whole point of Bullseye is 3-in, 1-out — not do-everything.
7	Build 50-lead sheet, pitch by hand	Warm converts fastest; cold fills the rest. Track the funnel, not vibes.
8	Credibility layer: BFS + reviews	The trust gap Charlie named. Runs alongside everything.
9	Free CRO report as the hook	Not a separate channel — the ammunition for steps 7 and 8.

1 · Define the ICP

Do not sell to “Shopify stores.” Define an Ideal Customer Profile tight enough that you can name 50 real stores against it. A usable segment is **industry + store size + a specific pain + a trigger event**. The trigger is what turns a fit into a *now*.

Working draft (to be confirmed by the interviews in §2):

A **Shopify DTC store with real traffic and a flat or declining conversion rate**, who today eyeballs GA4 or simply guesses at what’s broken — with no dev resource on call. The **trigger**: they just ran a paid-traffic push, hired a first growth/ops hire, or plateaued after a redesign.

Vocabulary rule. This buyer searches “traffic but no sales” and “leaking funnel” — never “CRO” or “AI expert.” Every line of copy, outreach and listing must match their words, not our internal ones.

Buying committee — collapsed. At this store size the owner or ops lead is economic buyer, champion and end user at once, with no procurement blocker. That is precisely why a fast, self-serve + founder-pitch motion fits.

2 · Interview — start with the warm 18

This is the correction Charlie pushed and it is the right one: **spend time interviewing your ideal customer**, don’t ask peers to grade your listing. Your own installed stores are the fastest, warmest source of truth — they have already used the product.

What each interview must extract:

- The **exact words** they use for the problem (these become your headline and outreach openers).
- What made them install — and what stopped them short of paying.
- Which single output they found genuinely useful (heatmap? replay? SmartNudge? a finding?).
- What they’d have to see to justify \$14–\$29/month — the real willingness-to-pay signal.

Reach every genuinely-active free store; two are already at their session cap, which is the easiest possible reason to open a conversation. Ranked by real usage:

STORE	SESSIONS / CAP	WHY THEY'RE FIRST
mexican-oaxacan-silver-jewelry	1500 / 1500 (100%)	Capped — losing visibility right now
freshleafindia	1500 / 1500 (100%)	Capped — losing visibility right now
esselbath1	1191 / 1500 (79%)	About to cap
cofainc	489 / 1500	Real, steady usage
bd00fa-3	352 / 1500	Longest-installed real user (since Feb)
skyline-decor	338 / 1500	Real, steady usage
kitchener-printkia	322 / 1500	Real, steady usage
corporatedrivertrainingaustralia	170 / 1500	Real, steady usage

These 8 are also your first **paying** cohort — the pitch is “you’re at/near your free limit, here’s what upgrading unlocks,” not a cold sell.

3 · Pick ≤ 3 value points — from the interviews, not from us

Charlie's sharpest note: "choose up to 3 key points to focus on — do not make your customer see a table with a lot of stuff." Today the listing and site sprawl. After the interviews you will know which 2–3 outcomes merchants actually care about, in their language. Everything downstream — video, images, site, listing, outreach — leads with those, and nothing else.

The candidate spine (to be confirmed by §2):

See why visitors don't buy → get the specific fix → recover the lost sales. The one honest differentiator left after the apply-engine was parked is that **SmartNudge deploys the intervention** — it doesn't just report the leak. If the interviews confirm that lands, it becomes value point #1.

4 · Remake the promo video, images & website

Only now — with validated words in hand. Doing this first would repeat the exact mistake Charlie flagged: an unfocused message and a site that "looks like AI vibe code" and doesn't match the brand. Order of work:

- 1 **Listing rewrite first** (it's been drafted since May and never shipped): app name, keywords in merchant vocabulary, hero line, description, screenshots. No engineering needed.
- 2 **Promo video** — 60 seconds: install → first finding → SmartNudge deployed, on a real store.
- 3 **Promo images** — rebuilt around the ≤ 3 points; drop the feature-dump table.
- 4 **Website / product voice** — match the brand, kill the AI-generated feel, lead with the outcome for the specific store type, one proof element beneath it.

Reality check. At near-zero traffic, polishing the site produces zero paying customers directly — it only improves conversion of traffic you don't have yet. Do the minimum listing fix now; save the full design lift for after the first interviews and first payers fund it.

5 · Traction research — map the outer ring

Bullseye: brainstorm every channel (outer ring), run cheap tests on the promising few (middle ring), then pour everything into the one that works (inner ring). Two research jobs before you test:

- **List all 19 traction channels** and write one concrete DynoWeb idea per channel — specific, not "SEO" but "rank a 'Lucky Orange alternative' page aimed at Shopify merchants."
- **Study your own category on the App Store** — how do Clarity, MIDA, Lucky Orange and the newer entrants actually get installs? Reviews, keywords, Built-for-Shopify badge, free tier, category placement. Their traction sources are visible; copy the mechanics.

Default priority order for digital B2B: founder-led sales & outbound, partnerships/integrations, content + SEO, email, engineering-as-marketing — with one addition specific to us: **the Shopify App Store is itself a search engine**, and ranking inside it is a real channel, not a "rarely-primary" one.

6 · Test 3 channels — pick 1 (the Bullseye trap)

The discipline is **3-in, 1-out**. Do not run SEO + GEO + backlinks + free features + outbound all at once — that is the disease. Define the success metric *before* running; "signal" means qualified conversations, not raw traffic.

CHANNEL	CHEAP TEST	SUCCESS SIGNAL
Convert warm free users	Founding offer to the 8 active stores in \$2	≥ 3–5 upgrade to paid
Founder-led outbound	30–40 personalized touches/week to named accounts	≥ 3 replies / 1 call per 30
Comparison-page SEO	Publish 1 “[competitor] alternative” page, then leave it	Ranks / 1 inbound in 2–3 wks

Deferred on purpose: SEO / GEO / backlinks are correct long-term but compound over months — they will produce ~0 paying users inside 60 days. They are the month-3+ layer, not this sprint’s engine. Say so out loud so they don’t quietly eat the hours.

7 · Build the 50-lead sheet & pitch by hand

From the confirmed ICP, list 50 named stores that fit — real names, pulled from App Store reviewers, stores matching the trigger, and lookalikes of whoever converted warm. Pitch each by hand and maintain the sheet. The columns are the funnel stages, because **the funnel math — not gut feel — tells you whether the plan can hit its number.**

Lead-sheet columns:

Store · contact · trigger · free-CRO-report sent? · touch 1/2/3 · replied? · call booked? · trialing? · **paid?** · notes

The honest funnel math (the plan’s Day 3 — do not skip it)

You set 50 paying in 60 days. Work it backwards before committing to it:

SOURCE	REALISTIC CONVERSION	PAYING
Warm — 8 active free stores	Founding offer, already getting value, some capped	~3–5
Cold — 50-person hand list	reply 15% × call 30% × close 30% ≈ 1.3% lead→paid	~1
Cold — to net ~15 paying	would need ~1,100 leads touched (impossible by hand)	—

Realistic 60-day ceiling: ~8–15 paying, not 50.

A 50-person by-hand list and a 50-paying goal are mathematically incompatible — one has to change. The recommendation: keep **50 as the north star you scale into**, and make the **real 60-day win**: first 8–15 paying + 5–10 App Store reviews + BFS badge applied + a pitch that converts at a *known* rate. Hit that and 50 becomes a volume problem you already know how to solve. Set 50-in-60 as the hard number and 15 will feel like failure — when 15 first-ever paying customers from zero is a real win.

8 · Credibility layer — BFS badge + reviews

Charlie named the real blocker: the app **lacks credibility — not Built-for-Shopify, and almost no reviews**. In B2B, trust and proof dominate; this is not a nice-to-have, it is the #4 principle of the whole plan. Runs in parallel with everything else.

- **Apply for the Built-for-Shopify badge in Week 1** — it lifts App Store ranking 30–50% and is the credibility stamp merchants look for.
- **Founding-Merchant offer:** first 15 payers get Growth (\$14/mo) at 30% off for 6 months in exchange for a written review once they’ve seen a real result. This closes faster than “try it and maybe upgrade,” builds the

review count, and produces case-study material for the next round of outreach. Cap at 15 — the scarcity is real.

- Every review and every case study warms the next deal — that is the **proof loop**, and it is the one growth loop already within reach.

9 · The free CRO report — the hook, not a fifth channel

This is the standout among the marketing ideas, but its power is that it is **not a separate workstream**. “I ran a free CRO report on your store — here’s the top leak costing you sales” is the outreach opener that actually earns replies, and it is the proof that supports the founding offer. Fold it into steps 7 and 8; do not run it as its own channel.

Engineering-as-marketing. A free, per-store CRO report is a classic free-tool traction play: it gives value before asking for anything, it is personalized (so it can’t be ignored like a template), and it demonstrates the product’s core promise in the first touch. It is the single best piece of ammunition you have — point it at the warm list and the named 50.

The 90-day cadence at a glance

WHEN	FOCUS	TARGET (cumulative paying)
Week 1	Ship the overdue listing fix · apply for BFS · start warm interviews · build the named-50 list	0 — groundwork
Weeks 2–4	Founding offer to the 8 warm stores · 30–40 outbound touches/week with the free CRO report as hook · track the funnel	8–10
Weeks 5–8	Kill the weaker lever, double the winner (Bullseye) · watch for review-driven App Store lift	18–20
Weeks 9–12	Referral ask baked into onboarding · finish the named list · compound toward the north star	30 — and climbing to 50

What has to be true — said plainly

This is aggressive for a motion with zero proven paid conversions. It only compounds if, inside the first three weeks, at least one is true: (a) the warm free-tier users convert meaningfully better than cold outbound, or (b) the Founding-Merchant discount is steep enough to overcome “nobody has paid for this yet.” If Week-2’s scorecard shows neither converting, the honest move is to revisit the offer or the number — not to run the same outreach at higher volume and hope.