DynoWeb Distribution Plan

Companion docs: dynoweb-icp.md — the canonical ICP (read first; every list, keyword and channel bet is gated by it) · dynoweb-offer.md — the canonical offer (positioning, entry offer, pricing, objection kill; the *what we sell* behind the *how we reach*).

Updated: 2026-07-17 — materially revised after the ICP research (§0); offer doc added (see §8–9).

0 · What the ICP research changed

The ICP work produced one finding that rewrites this plan's assumptions:

> **Zero of our ~18 real installed stores sit inside the ICP.** Exactly one clears the 5,000 sessions/month detectability floor — and it's a jewelry store, the one vertical we should avoid. The base skews heavily INR. Most stores are 5–10x below the traffic level at which our instruments can honestly see anything.

Our 0 paying customers is not a sales problem or a pricing problem. It is an acquisition-targeting problem. We acquired stores structurally incapable of getting recurring value: they install, the heatmap shows noise, there's nothing to act on, they never pay.

Two revisions follow:

1. **The warm-18 is downgraded from "fastest revenue" to "best interviews."** Earlier drafts of this plan assumed 3–5 paying conversions from warm users. On the evidence, that's optimistic — selling a paid plan to a below-floor store violates philosophy P8 and earns a refund plus a 1-star review. Interview them. Don't count on their money. (One exception: `mexican-oaxacan-silver-jewelry` is above-floor *and* capped — a real upgrade conversation.)
2. **The named-50 list is a hunting spec, not a description of our base.** It must be built from stores we don't have yet, gated by the ICP.

1 · Starting point (verified, not estimated)

Fact	Value	Source
Real paying customers	**0**	`Billing` — 4 "paid" rows are a test store, a seeded showcase, a dev store, and a partner
Real installed stores	~18 (of 25 rows)	`Shop` where `uninstalledAt IS NULL`, minus `isTestStore`
Stores inside the ICP	**0**	`TrafficDailyRollup`, 30d
Install rate	~1/week, stalled 6 weeks	`Shop.installedAt` by week
Listing rewrite	**Drafted May 20, never shipped**	`shopify.app.toml` still `name = "DynoWeb"`

Fact	Value	Source
Time budget	15–20 hrs/week	Confirmed

Treat month 1 as validation, not volume. Zero self-serve trial→paid conversions have ever happened. We are proving a paid motion exists — not scaling one.

2 · The plan in order

The spine: **you cannot write the message, shoot the video, or pitch the list until you have listened to merchants.** Everything before the interviews is preparation; everything after is execution.

#	Step	Why it comes here
1	Define the ICP	■ **Done** — dynoweb-icp.md
2	Interview — warm 18 first	The customers hand you the words. Before any asset work.
3	Pick ≤3 value points	No focused value prop today. The interviews decide the 3.
4	Remake video, images, site	Only now — around validated words, not guesses.
5	Bullseye traction research	Map all 19 channels + how App Store competitors get installs.
6	Test 3 channels, pick 1	3-in, 1-out. Not do-everything.
7	Build the 50-lead sheet, pitch by hand	ICP-gated. Track the funnel, not vibes.
8	Credibility: BFS badge + reviews	The trust gap. Runs alongside everything.
9	Free CRO report as the hook	Not a channel — the ammunition for 7 and 8.

3 · Interview the warm 18

This is the correction Charlie Ngo pushed and it's right: **interview your ideal customer; don't ask peers to grade your listing.**

Reframed goal: these interviews are now primarily **diagnostic**, not sales. The single most valuable question we can answer is *what did these below-floor stores think they were installing?* Their answer tells us what our listing currently promises, and to whom it's wrongly aimed.

What each interview must extract:

- The **exact words** they use for the problem → these become headline and outreach openers.
- What made them install — and what stopped them short of paying.
- Which single output they found genuinely useful (heatmap? replay? a finding? a nudge?).
- Whether they ever saw *anything actionable* — the below-floor hypothesis is that they did not.

Priority order (live traffic, 30d — not the capped Billing counter):

Store	Sessions/ 30d	CR	Why
mexican-oaxacan-silver-jewelry	5,511	0.20%	Above floor + capped. **The one real upgrade conversation.**

Store	Sessions/ 30d	CR	Why
freshleafindia	2,540	1.18%	Best revenue ($11.3k/30d). Below floor, wrong geo — but a real business.
esselbath1	2,050	**0.00%**	Zero conversions on 2k sessions. Either broken or a genuine find.
cofainc	1,481	0.00%	Real usage, no conversions
bd00fa-3	940	0.00%	Longest-installed real user (since Feb)
kitchener-printkia	687	0.87%	Converting — ask what's working
skyline-decor	681	0.15%	Real, steady
corporatedrivertrainingaustralia	390	0.77%	Real, steady

4 · Pick ≤3 value points — from interviews, not from us

Charlie's sharpest note: *"choose up to 3 key points to focus on — do not make your customer see a table with a lot of stuff."*

Candidate spine, to be confirmed by §3:

> **See why visitors don't buy → get the specific fix → recover the lost sales.**

The one honest differentiator after the apply-engine was parked is that **SmartNudge deploys the intervention** — it doesn't just report the leak.

> ■■ **Moat honesty check.** Wisepops and OptiMonk already ship popups with revenue attribution and A/B testing. "Deployed intervention + attribution" is **not** itself unique. What's plausibly unique is the **loop** — diagnose from replay/error data → deploy the nudge → attribute the money, in one app. Pressure-test this with a merchant before it becomes load-bearing marketing.

Vocabulary is binding (full table in the ICP doc §6): use *"lost sales"*, *"traffic but no sales"*, *"drop off"*, *"stuck"*. **Never** *"CRO"*, *"leaking funnel"*, or *"AI CRO Expert"* — all three are verified vendor-speak that merchants never use.

5 · Remake the promo video, images & website

Only now, with validated words in hand. Doing this first repeats the exact mistake Charlie flagged: an unfocused message on a site that *"looks like a AI Vibe Code web."*

1. **Listing rewrite first** (drafted since May, never shipped): app name, keywords in merchant vocabulary, hero line, description, screenshots. No engineering needed.
2. **Promo video** — 60s: install → first finding → SmartNudge deployed, on a real store.
3. **Promo images** — rebuilt around the ≤3 points. Drop the feature-dump table.
4. **Website / product voice** — match the brand, kill the AI-generated feel, lead with the outcome for a specific store type, one proof element beneath it.

> **Reality check.** At near-zero traffic, polishing the site produces zero paying customers directly — it only improves conversion of traffic we don't have yet. Do the minimum listing fix now; let the first payers fund the full design lift.

6 · Bullseye — map the outer ring, test 3, pick 1

Research first:

- List all **19 traction channels**, one concrete DynoWeb idea each — specific, not "SEO" but "rank a *Lucky Orange alternative* page aimed at Shopify merchants."
- **Study our own category on the App Store** — how do Clarity, MIDA, Lucky Orange and Propel actually get installs? Reviews, keywords, BFS badge, free tier, category placement. Their traction mechanics are visible. Copy them.

One addition specific to us: the **Shopify App Store is itself a search engine**, and ranking inside it is a real channel — not the "rarely primary" the generic framework calls it.

Then test 3 — and only 3:

Channel	Cheap test	Success signal
Founder-led outbound	30–40 personalized touches/week to ICP-gated named accounts	≥3 replies / 1 call per 30
App Store listing + reviews	Ship the rewrite, push for reviews, apply for BFS	Install rate or ranking movement in 2–3 wks
Comparison-page SEO	Publish 1 "[competitor] alternative" page, then leave it	Ranks / 1 qualified inbound in 2–3 wks

> **Deferred on purpose:** SEO / GEO / backlinks are correct long-term but compound over *months* — they will produce ~0 paying users inside 60 days. They're the month-3+ layer, not this sprint's engine. Say it out loud so they don't quietly eat the hours.

---

7 · The 50-lead sheet — ICP-gated, pitched by hand

Build from the ICP's gates (§2 of that doc). **A lead passes all six gates or it doesn't go on the sheet.** A big list of unqualified stores is worse than a small qualified one — it manufactures activity and hides that we're not reaching the right people.

Sheet columns: `store · contact · trigger · sessions/mo (est) · vertical · geo · free-CRO-report sent? · touch 1/2/3 · replied? · call booked? · trialing? · paid? · notes`

The honest funnel math

Source	Realistic conversion	Paying
Warm — 8 active free stores	**Mostly below floor.** Only 1 is a real upgrade conversation	**~1–2**
Cold — 50-person hand list	reply 15% × call 30% × close 30% ≈ **1.3% lead→paid**	~1
Cold — to net ~15 paying	would need **~1,100 leads touched** (not possible by hand)	—

> **Realistic 60-day ceiling: ~5–12 paying — and it was ~8–15 before the ICP finding.**
>
> A 50-person by-hand list and a 50-paying goal are mathematically incompatible; one has to change. **Recommendation:** keep **50 as the north star we scale into**, and make the **real 60-day win**: first 5–12 paying + 5–10 App Store reviews + BFS applied + a pitch converting at a *known* rate. Hit that and 50 becomes a volume problem we already know how to solve. Set 50-in-60 as the hard number and 10 will feel like failure — when 10 first-ever paying customers, from zero, from a correctly-targeted ICP, is the actual win.

---

8 · Credibility — BFS badge + reviews

Charlie named the real blocker: the app **lacks credibility — not Built-for-Shopify, and almost no reviews.** In B2B, trust and proof dominate. This is not a nice-to-have.

- **Apply for the Built-for-Shopify badge in Week 1** — lifts App Store ranking 30–50% and is the credibility stamp merchants look for.
- **Founding-Merchant offer:** first 15 payers get Growth ($14/mo) at 30% off for 6 months in exchange for a written review once they've seen a real result. Closes faster than "try it and maybe upgrade," builds review count, produces case-study material. Cap at 15 — the scarcity is real. → **The full offer mechanics** (the outcome-gated "Free Leak Report + First Fix Free" entry, and the $149 done-for-you tier that breaks the self-serve ceiling) live in dynoweb-offer.md — the founding-merchant deal is one lever inside it.
- Every review and case study warms the next deal. That's the **proof loop** — the one growth loop within reach.

The wedge against free is not features. Clarity's 1-star reviews are unanimous: *"would not let me log in"*, *"no support… zero response"*, *"very difficult for a shopkeeper to grasp"*, *"data basically unusable… all garbled."* **Not one says "I want more features."** Setup, comprehensibility, support, and Shopify-native context are the opening.

9 · The free CRO report — the hook, not a fifth channel

"I ran a free CRO report on your store — here's the top leak costing you sales" is the outreach opener that earns replies, and the proof that supports the founding offer. It gives value before asking, it's personalized (so it can't be ignored like a template), and it demonstrates the core promise in the first touch.

Fold it into §7 and §8. Do not run it as its own workstream.

10 · The 90-day cadence

When	Focus	Target (cumulative paying)
Week 1	Ship the overdue listing fix · apply for BFS · start warm interviews · build the ICP-gated named-50	0 — groundwork
Weeks 2–4	Founding offer where honest · 30–40 outbound touches/week with the free CRO report as hook · track the funnel	3–5
Weeks 5–8	Kill the weaker lever, double the winner · watch for review-driven App Store lift	10–14
Weeks 9–12	Referral ask in onboarding · finish the named list · compound toward the north star	20–30

11 · What has to be true — said plainly

This is aggressive for a motion with **zero proven paid conversions and zero ICP-fit stores in the base.** It compounds only if, inside the first three weeks, at least one is true:

- **(a)** ICP-gated cold outbound converts meaningfully better than our historical install→paid rate (which is currently 0% — a low bar, but it has never been cleared), or
- **(b)** the Founding-Merchant discount is steep enough to overcome *"nobody has paid for this yet."*

If Week 2's scorecard shows neither converting, the honest move is to revisit the offer, the ICP, or the number — **not** to run the same outreach at higher volume and hope.

The deeper bet: that the ICP is right and our base was simply wrong. If ICP-gated stores also fail to convert, the problem is the product's value at this price, not the targeting — and that's a different, larger conversation worth having early rather than after 1,000 cold emails.